

Assignment-3

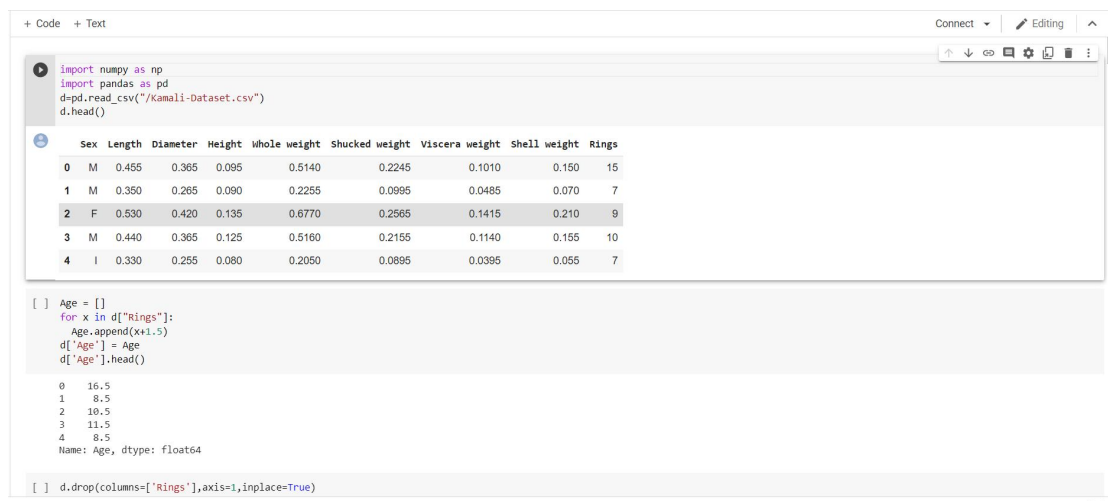
Statistical Machine Learning Approaches To Liver Disease Prediction

Student Name	Kamali R
Student Roll no	621319104019
Maximum Marks	2 Marks

Abalone Age Prediction:

Download the dataset

Load the dataset into the tool.



```
+ Code + Text
```

```
import numpy as np
import pandas as pd
d=pd.read_csv("/Kamali-Dataset.csv")
d.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.485	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

```
[ ] Age = []
for x in d[["Rings"]]:
    Age.append(x+1.5)
d['Age'] = Age
d['Age'].head()

0    16.5
1     8.5
2    10.5
3    11.5
4     8.5
Name: Age, dtype: float64

[ ] d.drop(columns=["Rings"],axis=1,inplace=True)
```

Univariate Analysis

Bi-Variate Analysis

Multi-Variate Analysis

+ Code + Text


Connect Editing ^

```
d.drop(columns=['Rings'],axis=1,inplace=True)
d.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

```
[ ] import seaborn as sns
sns.displot(d["Age"])
```

<seaborn.axisgrid.FacetGrid at 0x7f11d47f38d0>




+ Code + Text

Connect Editing ^

```
sns.histplot(x=d['Age'])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f11c5a9d410>



```
[ ] sns.boxplot(x=d['Age'])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f11c553bd10>




+ Code + Text

Connect ^


```
sns.countplot(x=d['Age'])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f11c552f990>




```
[ ] sns.barplot(x=d['Height'],y=d['Age'])
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f11c5353d90>



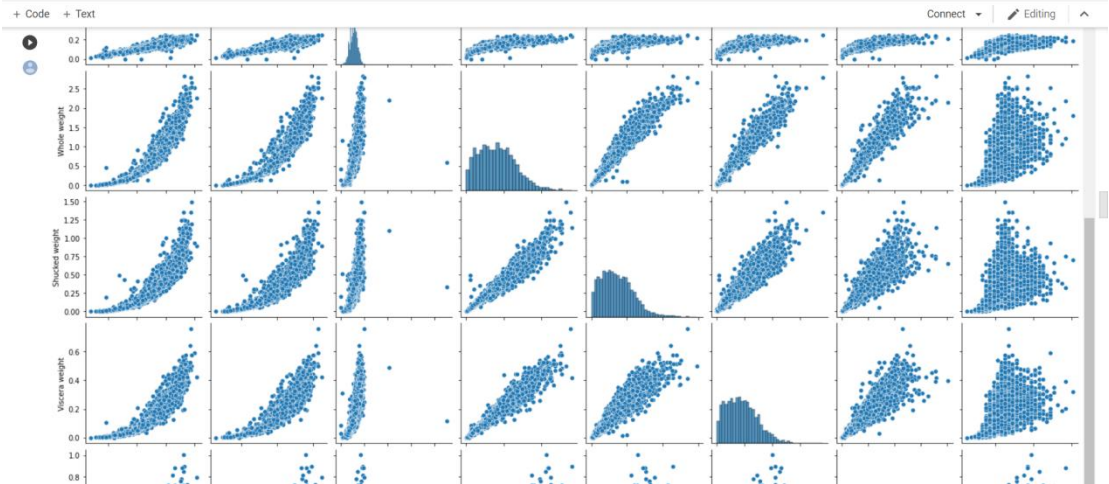

```
+ Code + Text
sns.lineplot(x=d['Age'],y=d['Height'])
```




```
[ ] sns.scatterplot(x=d['Age'],y=d['Height'])
```



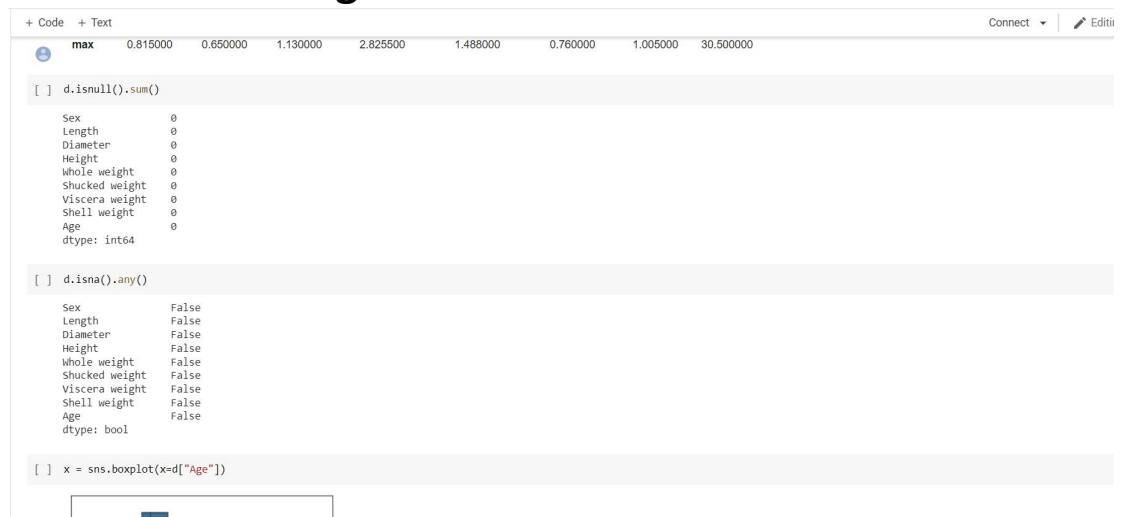
```
[ ] sns.regplot(x=d['Age'],y=d['Height'])
```



Perform descriptive statistics on the dataset.



Check for Missing values and deal with them.



```
+ Code + Text Connect Editing ^

[ ] x = np.where(d['Age']>57,39, d['Age'])
x
array([16.5,  8.5, 10.5, ..., 10.5, 11.5, 13.5])

sns.boxplot(x=x)
<matplotlib.axes._subplots.AxesSubplot at 0x7f11c17ff310>


[ ] import warnings
warnings.filterwarnings('ignore')
pd.Categorical(d['whole weight'])

[0.5140, 0.2255, 0.6770, 0.5160, 0.2050, ..., 0.8870, 0.9660, 1.1760, 1.0945, 1.9485]
Length: 4177
Categories (2429, float64): [0.0020, 0.0080, 0.0105, 0.0130, ..., 2.5550, 2.6570, 2.7795, 2.8255]

[ ] pd.get_dummies(d['Height']).head()
```

Split the data into dependent and independent variables.

Check for Categorical columns and perform encoding.

```
+ Code + Text Connect Editing ^

pd.get_dummies(d['Height']).head()
0.000 0.010 0.015 0.020 0.025 0.030 0.035 0.040 0.045 0.050 ... 0.210 0.215 0.220 0.225 0.230 0.235 0.240 0.250 0.515 1.130
0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0
1 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0
2 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0
3 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0
4 0 0 0 0 0 0 0 0 0 0 0 ... 0 0 0 0 0 0 0 0 0 0
5 rows x 51 columns

[ ] pd.get_dummies(d).head()
Length Diameter Height whole weight Shucked weight Viscera weight Shell weight Age Sex_F Sex_I Sex_M
0 0.455 0.365 0.095 0.5140 0.2245 0.1010 0.150 16.5 0 0 1
1 0.350 0.265 0.090 0.2255 0.0995 0.0485 0.070 8.5 0 0 1
2 0.530 0.420 0.135 0.6770 0.2565 0.1415 0.210 10.5 1 0 0
3 0.440 0.365 0.125 0.5160 0.2155 0.1140 0.155 11.5 0 0 1
4 0.330 0.255 0.080 0.2050 0.0895 0.0395 0.055 8.5 0 1 0

[ ] X = d.iloc[:, :-1].values
X
```

Scale the independent variables

```
+ Code + Text Connect Editing ^

X = d.iloc[:, :-1].values
X
array([[ 'M', 0.455, 0.365, ..., 0.2245, 0.101, 0.15],
       [ 'M', 0.35, 0.265, ..., 0.0995, 0.0485, 0.07],
       [ 'F', 0.53, 0.42, ..., 0.2565, 0.1415, 0.21],
       ...,
       [ 'M', 0.6, 0.475, ..., 0.5255, 0.2875, 0.308],
       [ 'F', 0.625, 0.485, ..., 0.531, 0.261, 0.296],
       [ 'M', 0.71, 0.555, ..., 0.9455, 0.3765, 0.495]])

[ ] Y = d.iloc[:, -1].values
Y
array([16.5,  8.5, 10.5, ..., 10.5, 11.5, 13.5])

[ ] from sklearn.preprocessing import scale
x = scale(d['Viscera weight'])
x
array([-0.72621157, -1.20522124, -0.35668983, ...,  0.97541324,
        0.73362741,  1.78744868])

[ ] x = d.iloc[:, 1:7]
x
Length Diameter Height whole weight Shucked weight Viscera weight
0 0.455 0.365 0.095 0.5140 0.2245 0.1010
1 0.350 0.265 0.090 0.2255 0.0995 0.0485
```

+ Code
+ Text

Connect
Editing

[]	1	0.350	0.265	0.090	0.2255	0.0995	0.0485
	2	0.530	0.420	0.135	0.6770	0.2565	0.1415
	3	0.440	0.365	0.125	0.5160	0.2155	0.1140
	4	0.330	0.255	0.080	0.2050	0.0895	0.0395

	4172	0.565	0.450	0.165	0.8870	0.3700	0.2390
	4173	0.590	0.440	0.135	0.9660	0.4390	0.2145
	4174	0.600	0.475	0.205	1.1760	0.5255	0.2875
	4175	0.625	0.485	0.150	1.0945	0.5310	0.2610
	4176	0.710	0.555	0.195	1.9485	0.9455	0.3765

4177 rows × 6 columns

```

y = d.iloc[:, -1]
y
0      16.5
1       8.5
2      10.5
3      11.5
4       8.5
...
4172    12.5
4173    11.5
4174    10.5
4175    11.5
4176    13.5

```

- Split the data into training and testing**
- Build the Model**
- Train the Model**
- Test the Model**

```
+ Code + Text
4175 11.5
4176 13.5
Name: Age, Length: 4177, dtype: float64

[ ] from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state =42)
x_train.shape

(3132, 6)

[ ] y_test.shape

(1045,)

[ ] x_train.head()

Length Diameter Height whole weight Shucked weight Viscera weight
3823 0.615 0.455 0.135 1.0590 0.4735 0.2630
3956 0.515 0.395 0.140 0.6860 0.2810 0.1255
3623 0.660 0.530 0.175 1.5830 0.7395 0.3505
0 0.455 0.365 0.095 0.5140 0.2245 0.1010
2183 0.495 0.400 0.155 0.8085 0.2345 0.1155

[ ] x_test.head()

Length Diameter Height whole weight Shucked weight Viscera weight
```

```
+ Code + Text
```

```
x_test.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
866	0.605	0.455	0.160	1.1035	0.4210	0.3015
1483	0.590	0.440	0.150	0.8725	0.3870	0.2150
599	0.560	0.445	0.195	0.9810	0.3050	0.2245
1702	0.635	0.490	0.170	1.2615	0.5385	0.2665
670	0.475	0.385	0.145	0.6175	0.2350	0.1080

```
[ ] y_train.head()
```

```
3823    10.5
3956    13.5
3623    11.5
0        16.5
2183     7.5
Name: Age, dtype: float64
```

```
[ ] y_test.head()
```

```
866      10.5
1483     9.5
599      17.5
1702     10.5
670      15.5
Name: Age, dtype: float64
```

Measure the performance using Metrics.

```
+ Code + Text Connect Ed

from sklearn.linear_model import LinearRegression
model=LinearRegression()
model.fit(x_train,y_train)

LinearRegression()

[ ] Y_predict_train = model.predict(x_train)
Y_predict_train

array([11.25888828, 11.95379472, 12.33692259, ..., 11.12903068,
       10.71152746, 11.59516371])

[ ] y_predict = model.predict(x_test)
y_predict

array([13.0478407 , 11.43166184, 15.59825921, ..., 13.69440346,
       11.79279231, 10.83037939])

[ ] import math
from sklearn.metrics import mean_squared_error
print(mean_squared_error(y_test, y_predict))
print(math.sqrt(mean_squared_error(y_test, y_predict)))

4.862459933051861
2.2050986220692854
```